

CerbaSeal

Core v0.1.0

Commercialization, Pilot Readiness & Market Positioning Dossier

A neutral partner-ready dossier combining technical reference, guided operations, pilot model, market evidence, pricing framework, and commercialization roadmap.

VERSION

0.1.0

STATUS

Pilot and Partner Evaluation

DATE

June 2026

AUTHOR

Jesse Lamont / Lamont Labs

Table of Contents

This dossier is structured for a partner or commercialization lead who needs to understand what CerbaSeal is, how it can be piloted, how it can be delivered repeatably, and how a pricing model can be discussed using current market evidence.

Section	Title
1	Executive Summary
2	Commercial Context and Strategic Thesis
3	System Overview: What CerbaSeal Is
4	Technical Architecture in Business Terms
5	Governance Model and Invariant Framework
6	Configuration, Policy, and Workflow Mapping
7	Audit, Evidence, and Replay Model
8	Deployment and Guided Operations Model
9	Pilot Delivery Model
10	Partner Delivery and Founder Independence
11	Market Category and Comparable Benchmarks
12	Pricing Evidence and Commercial Model
13	Partner Economics and Channel Strategy
14	Adoption Roadmap and Build Priorities
15	Risk Register and Mitigation Plan
16	Recommended Next Steps
17	Partner Discussion Brief
18	Pilot Commercial Structure
Appendix A	Reason Codes
Appendix B	Deployment Checklist
Appendix C	Source Notes

Source basis. This document is derived from the CerbaSeal-Core v0.1.0 technical reference manual supplied by Lamont Labs, plus market-pricing and commercialization research gathered from public AWS Marketplace, EU AI Act, AI governance, policy enforcement, middleware, and channel-partner sources. It is written for partner evaluation and pilot planning, not as a final legal, security, or compliance opinion.

1. Executive Summary

CerbaSeal-Core is a deterministic governance enforcement layer for organizations that need a controlled boundary between decision generation and consequential execution. It is intended for AI-assisted and automated workflows where a recommendation, escalation, approval, account action, transaction decision, or tool execution should not proceed until defined governance requirements have been evaluated.

The platform does not replace the client workflow engine, case management system, identity provider, compliance platform, or AI model. It provides the enforcement point those systems can call before an action proceeds. The output is a consistent ALLOW, HOLD, or REJECT decision supported by traceable inputs, invariant checks, reason codes, audit events, and exportable evidence.

Positioning statement: CerbaSeal should be evaluated as governance enforcement infrastructure: a deterministic control layer that integrates with existing systems to enforce approval, provenance, logging, trust, and policy requirements before execution.

Dimension	Current Position
Technical maturity	Working enforcement core with 391 passing tests across 17 test files and 16 repository governance checks documented in the current reference materials.
Pilot posture	Best framed for controlled pilots, partner evaluation, customer-hosted deployment exercises, and workflow-specific validation.
Commercial posture	Most comparable to runtime policy enforcement, AI governance infrastructure, and enterprise middleware rather than simple compliance automation.
Delivery posture	Guided independence: clients retain control of environment and workflow; partners and Lamont Labs provide deployment guidance, mapping support, validation, evidence review, and defined escalation.
Pricing posture	Pilot and annual licensing should be based on governance infrastructure, deployment support, evidence value, and partner delivery economics - not on commodity SaaS seat pricing.

The core commercial question is no longer simply whether CerbaSeal works. The evidence now supports technical credibility. The next question is whether CerbaSeal can be delivered repeatedly with bounded support, partner enablement, clear pricing, and a customer experience that does not depend on the founder being present for every operational detail.

2. Commercial Context and Strategic Thesis

Enterprise buyers do not buy governance infrastructure because it is technically elegant. They buy it because it reduces control ambiguity around important workflows. CerbaSeal creates an explicit authorization boundary between recommendation and execution, which is increasingly relevant as organizations use AI systems, automation layers, and agentic workflows near operational decision points.

For partner discussions, CerbaSeal should be described as an execution governance layer. The simplest commercial framing is: AI and automation can propose actions; CerbaSeal determines whether those actions are authorized to proceed.

Buyer-relevant outcomes

Outcome	What CerbaSeal contributes
Operational clarity	Every request returns ALLOW, HOLD, or REJECT with reason codes and checked governance controls.
Audit readiness	Evidence bundles and audit-chain verification support internal review, pilot evidence packages, and governance analysis.
Human oversight	AI proposals remain non-authoritative and must satisfy approval and policy requirements before execution.
Deployment control	The system can run in customer-controlled environments using Docker Compose or direct Node.js paths.
Partner scalability	A guided delivery framework enables partners to qualify, deploy, configure, verify, and hand over pilots with defined escalation boundaries.

The strongest near-term buyer profile is a regulated or control-sensitive organization with one narrow workflow where approval, traceability, and evidence matter. Examples include fraud triage, transaction escalation, account hold recommendation, KYC escalation, insurance claim review, or AI agent tool execution.

3. System Overview: What CerbaSeal Is

CerbaSeal-Core is organized around a simple operational contract: client systems construct a GovernedRequest, the gate evaluates that request against invariant and policy requirements, and the result determines whether execution may proceed, pause for approval, or stop.

Object	Role in the system
GovernedRequest	The complete structured input carrying workflow class, actor identity, actor authority, proposed action, proposal metadata, policy reference, provenance, approval artifact, logging readiness, control status, trust state, and timestamp.
ExecutionGateService	The deterministic evaluation component that returns GateResult.
GateResult	The output object containing DecisionEnvelope plus either ReleaseAuthorization or BlockedActionRecord.
DecisionEnvelope	The decision record containing final state, trace, evidence reference, and issuance information.
ReleaseAuthorization	Issued only when finalState is ALLOW. This is the record downstream systems can use to proceed.
BlockedActionRecord	Issued when finalState is HOLD or REJECT. This documents why an action did not proceed.
EvidenceBundle	The assembled package of request, decision, authorization or block record, and relevant audit entries.

Decision states

State	Meaning	Operational interpretation
ALLOW	All required checks passed.	Downstream execution may proceed according to the client workflow.
HOLD	The request is structurally valid, but required approval is absent.	The workflow should pause until approval evidence is provided or the request is resubmitted.
REJECT	An invariant or policy rule failed.	The current request should not proceed. A corrected request is required before another evaluation.

4. Technical Architecture in Business Terms

The source reference manual defines CerbaSeal as a seven-layer architecture. Each layer has a specific responsibility, and together they create a structured path from request submission to decision evidence.

Layer	Component	Business purpose
1 - Request	GovernedRequest	Captures the proposed action and all governance-relevant context in one structured object.
2 - Trust	TrustState	Represents whether the requesting actor is trusted at evaluation time.
3 - Provenance	ProvenanceRef	Connects the proposal to model version, rule set version, and source hash.
4 - Gate	ExecutionGateService	Evaluates baseline controls and returns a deterministic decision.
5 - Policy	CerbaSealPolicy	Applies client-specific actor mappings, approval chains, workflow rules, and action policies.
6 - Audit	AppendOnlyLogService / FileBackedAppendOnlyLogService	Records append-only audit events and verifies hash-chain continuity.
7 - Evidence	EvidenceBundleService	Assembles proof material for review, handoff, pilot closeout, or investigation.

Evaluation lifecycle

Stage	Primary owner	Summary
Construct	Client integration	Build a GovernedRequest from AI output, identity context, workflow metadata, and control status.
Submit	Client integration	Call gate.evaluate(request). Evaluation is synchronous.
Schema and boundary checks	Gate	Validate request shape, known action class, actor authority class, and proposal/action match.
Policy stages	Gate	Resolve actor mappings, blocked action policies, effective approval requirements, and approval-chain authority.
Invariant checks	Gate	Evaluate policy pack, provenance, logging readiness, AI boundary, prohibited use, controls, trust state, and approval state.
Decision	Gate	Return ALLOW, HOLD, or REJECT with envelope and authorization or block record.
Evidence	Caller / evidence service	Append audit events and assemble evidence bundle for review.
Export	Partner / operator	Run proof export and verification to confirm evidence package integrity.

5. Governance Model and Invariant Framework

The invariant framework provides the baseline governance contract. Invariants are unconditional in v0.1.0: client policy can add requirements but cannot remove core requirements. This distinction matters commercially because it separates organization-specific configuration from non-negotiable execution controls.

Code	Control name	Plain-language meaning
INV-01	No policy pack, no execution	policyPackRef must be present.
INV-02	No provenance, no action	provenanceRef must include model version, rule set version, and source hash.
INV-03	Required approval, no release	Required approval must be present and valid; otherwise HOLD or REJECT.
INV-04	No logging, no execution	loggingReady must be true before release can issue.
INV-05	AI non-authoritative	AI proposal source and AI actor authority cannot authorize execution.
INV-06	No bypass of execution gate	Evidence requires a gate-issued GateResult registered by evaluate().
INV-07	Immutable decision envelope	DecisionEnvelope carries immutable: true and is tied to gate-issued result flow.
INV-08	Sensitive controls must be current	Sensitive requests require critical controls valid and not stale.
INV-09	Trust state required	trustState.trusted must be true.
INV-10	Prohibited use blocks execution	prohibitedUse true returns REJECT regardless of other fields.
INV-11	Request schema and action class valid	Required schema values, authority class, and action class must be valid.
INV-12	Proposal/action match	proposedActionClass must match proposal.requestedActionClass.

Buyer-friendly interpretation

Risk pattern	CerbaSeal response
AI attempts to authorize itself	The gate rejects due to AI non-authoritativeness.
Approval absent where required	The gate returns HOLD, preserving human review without releasing execution.
Approval reused across requests	The gate rejects because approvalArtifact.forRequestId must match the request.
Missing provenance	The gate rejects because the proposal cannot be traced to model, rule set, and source hash.
Audit logging not ready	The gate rejects because release requires a ready audit path.
Controls stale on sensitive action	The gate rejects because stale or invalid controls cannot support sensitive release.

6. Configuration, Policy, and Workflow Mapping

The configuration system allows the core class sets to be extended and client-specific rules to be defined without modifying the enforcement gate. The current reference documents `cerbaseal.config.json` for class extensions and `cerbaseal.policy.json` for actor mappings, approval chains, workflow rules, and action policies.

Layer	Purpose	Example
Core classes	Always registered base system language.	system, ai, analyst, reviewer, manager, compliance_officer; fraud_triage; allow, hold, reject, escalate, account_hold.
Extended classes	Additional identifiers accepted in a client or pilot context.	risk_officer, senior_analyst, insurance_claim_review, kyc_escalation.
Actor mappings	Translate client role labels into authority classes used by the gate.	Head of Risk -> manager; Senior Fraud Analyst -> analyst.
Approval chains	Define which authority classes can approve a workflow.	fraud_triage -> manager, compliance_officer.
Workflow rules	Add workflow-level approval requirements.	kyc_escalation requires approval.
Action policies	Set action-specific behavior for a workflow.	allow requires_approval; reject auto_allow; account_hold blocked.

Configuration posture: The policy system is additive. It can make the gate more restrictive for a client workflow, but it is not designed to bypass or relax the invariant framework.

The major productization opportunity is to make configuration feel like governance setup rather than developer work. The highest-leverage future assets are a policy pack builder, workflow mapping wizard, and governance pack library that reduce JSON hand-authoring and make pilot setup repeatable.

7. Audit, Evidence, and Replay Model

Auditability is a central feature of CerbaSeal. The platform records decision-related events and supports verification of audit-chain continuity. Evidence bundles provide a structured review artifact for pilot closeout, support analysis, compliance discussions, and operational review.

Service	Deployment use	Notes
AppendOnlyLogService	Development, tests, and non-persistent demo use.	In-memory audit entries; returned values are cloned to avoid caller mutation.
FileBackedAppendOnlyLogService	Client deployments where audit entries must survive restart.	JSONL file; reads existing chain on startup; appendFileSync writes entries before returning.
verifyChain()	Operational health and audit integrity checks.	Returns true only when every SHA-256 link is valid.

Evidence bundle contents

Evidence element	Description
evidenceBundleId	Deterministic evidence identifier for the request.
GovernedRequest	Original request payload used for evaluation.
DecisionEnvelope	Final gate decision, checked invariants, reason codes, evidence reference, and issued timestamp.
ReleaseAuthorization or BlockedActionRecord	Authorization when ALLOW; block record when HOLD or REJECT.
eventChain	AuditLogEntry objects for the requestId.
createdAt	Evidence bundle creation timestamp.

The commercial value of the evidence model is that CerbaSeal does not merely decide. It creates a reviewable record of why a decision path proceeded, paused, or stopped. That record becomes useful in pilots, internal audits, operational reviews, and partner-led demonstrations.

8. Deployment and Guided Operations Model

CerbaSeal should not be positioned as either "self-service only" or "founder-managed forever." The strongest operating posture is guided independence: clients retain control of environment and workflow; partners and Lamont Labs provide a repeatable path for deployment, mapping, configuration, validation, evidence review, and escalation support.

Guided independence: the client is supported, the deployment is controlled, and routine operation is not dependent on a single founder or implementation engineer.

Responsibility area	Client role	Partner / Lamont Labs role
Environment ownership	Provide approved hosting environment, network controls, persistence volume, and operational monitoring.	Provide deployment guide, reference architecture, Docker Compose path, direct Node.js path, and validation steps.
Workflow definition	Identify business workflow, actor roles, approval path, and success criteria.	Guide workflow mapping into request, config, policy, and verification structures.
Identity and access	Operate identity provider, authentication, privileged access, and network-level API controls.	Document trust-state expectations and boundary assumptions; support integration planning.
Policy configuration	Review and approve how organizational roles and approval chains are represented.	Provide policy template, configuration guidance, validation scripts, and certification path.
Pilot execution	Run representative requests and review outcomes with stakeholders.	Facilitate setup, evidence capture, issue triage, and closeout.
Evidence review	Use evidence outputs for internal governance, audit preparation, and workflow review.	Provide export procedures, verification steps, and interpretation support during pilot.

Deployment options

Path	Recommended use	Summary
Docker Compose	Controlled pilots and partner deployments where containerization is allowed.	Containerized runtime, persistent audit volume, health check, and config/policy mount path.
Node.js Direct	Pilot configuration and source-level validation environments.	pnpm install, pnpm test, pnpm audit:repo, setup wizard, verify script, proof export, and direct server start.
Future package model	Post-v0.1 productization.	Published package, packaged templates, and lower-friction install path to reduce partner setup effort.

9. Pilot Delivery Model

A CerbaSeal pilot should be framed as a guided validation of one governed workflow, not a broad production rollout. The pilot is successful when the client and partner can show that a defined workflow can be mapped, evaluated, audited, and reviewed through CerbaSeal with clear operational responsibilities and evidence outputs.

Principle	Meaning
One workflow first	Begin with a single high-value workflow rather than broad platform adoption.
Guided implementation	Partner and Lamont Labs provide structure, templates, and support rather than leaving the client to interpret the system alone.
Client-controlled environment	Client retains control of deployment environment, identity systems, network controls, and business operations.
Defined support window	Support expectations are explicit, bounded, and aligned to pilot scope.
Evidence-first closeout	The pilot ends with evidence package, verification output, support log, and lessons learned.

Pilot timeline

Phase	Primary activities	Output
Discovery	Confirm workflow, actors, approvals, deployment environment, and success criteria.	Pilot scope and readiness confirmation.
Mapping	Translate workflow into GovernedRequest, class extensions, actor mappings, approval chains, and action policies.	Configuration and policy draft.
Deployment	Run Docker Compose or Node.js Direct path; configure logging and health checks.	Healthy system with verified audit chain.
Verification	Run ALLOW, HOLD, and REJECT scenarios against the mapped workflow.	Scenario results and verification log.
Pilot operation	Run representative requests; gather evidence; collect feedback and operational questions.	Observed outcomes and support notes.
Closeout	Export proof, verify proof, review evidence, record lessons, and define next steps.	Evidence package and pilot outcome summary.

The pilot should be priced and scoped as a commercial validation engagement. It includes discovery, mapping, deployment support, verification, evidence review, support window, and closeout. It is not a free trial and it is not a full production rollout.

10. Partner Delivery and Founder Independence

The partner framework is designed to make CerbaSeal deliverable through a repeatable, supportable process. The purpose is not to transfer all effort to partners or clients. The purpose is to equip partners with enough structure to lead routine delivery while Lamont Labs remains available for defined escalation cases.

Document	Primary audience	Purpose
00-OVERVIEW.md	All partners	Navigation, delivery model, certification levels, escalation boundaries.
01-sales-brief.md	Sales leads	ICP qualification, positioning, commercial model, opening objections.
02-technical-brief.md	Technical leads	Architecture, invariants, evidence chain, integration patterns.
03-deployment-guide.md	Engineers	Environment setup, wizard, audit, verification, evidence export.
04-pilot-guide.md	Pilot leads	1-day pilot path from discovery to evidence package.
05-support-guide.md	All partners	Common issues, root causes, diagnosis steps, and resolution path.
06-objection-handling.md	Sales leads	Enterprise objections with role-specific responses.
07-certification-framework.md	All partners	Certification competencies, practical exercises, pass/fail criteria.

Certification model

Level	Name	What it proves
1	Deploy	Partner can install, test, audit, verify, and export proof for the system.
2	Configure	Partner can author a policy file for a new client workflow and pass verification.
3	Lead pilots	Partner can run a pilot from kickoff through evidence closeout with appropriate escalation discipline.

Founder independence target

Current dependency area	Reduction mechanism	Target outcome
Workflow mapping questions	Workflow mapping workbook and future wizard.	Partner can map common workflows without Tier 3 support.
Policy authoring	Templates, policy pack builder, validation scripts.	Configuration becomes guided setup rather than source-level work.
Deployment troubleshooting	Docker Compose path, health endpoint, deployment checklist.	Partner resolves routine deployment issues.
Evidence interpretation	Evidence explorer, closeout template, training material.	Client and partner can review proof package without founder-led explanation.
Commercial scoping	Pilot tiers, pricing framework, and success criteria.	Partner can frame pilot offer consistently.

11. Market Category and Comparable Benchmarks

CerbaSeal should not be priced or positioned like lightweight compliance automation. The strongest analogs are runtime policy enforcement, AI governance control planes, decision governance infrastructure, and enterprise middleware. Those categories monetize control, auditability, integration, and operational assurance.

Category	Why it matters for CerbaSeal	Representative examples
Runtime policy enforcement	Closest structural analog: policy is evaluated in the path of execution.	Styra / OPA ecosystem, Axiomatics, PlainID, Permit.io.
AI governance control plane	Comparable buyer concern: auditability, oversight, controls, traceability, and policy outcomes.	Credo AI, Arthur AI, Fiddler AI.
Enterprise middleware	Comparable buying behavior: sits between systems and becomes operationally critical.	MuleSoft, Workato, Boomi, Tines.
Decision governance	Comparable workflow context: risk decisions, fraud decisions, transaction actions, approval chains.	FICO Decision Management, Pega Decisioning, Actimize, Feedzai.

Market evidence points to five- and six-figure annual contracts for adjacent infrastructure categories. Exact private enterprise contract values are often not public, so the most defensible framework uses evidence tiers: hard marketplace pricing, public custom-pricing statements, buyer-reported or third-party estimates, and adjacent category benchmarks.

Evidence point	Public source evidence	CerbaSeal implication
Arthur AI	AWS Marketplace and Arthur materials position Arthur as continuous evaluations and guardrails for AI, with public marketplace category availability.	AI governance infrastructure can be sold through marketplace and enterprise procurement channels.
Credo AI	AWS Marketplace states pricing depends on contract duration, contract terms, and usage.	Enterprise AI governance pricing is often private-offer and use-case based.
Styra	AWS Marketplace lists Styra DAS Enterprise at \$5,000/month for 16 systems and 100MM decisions/month.	Runtime policy enforcement has a concrete roughly \$60K/year public benchmark.
Workato	AWS Marketplace lists Enterprise workspace with 1,000,000 tasks and Enterprise Support at \$143,750/year.	Between-systems automation infrastructure supports six-figure annual pricing.
MuleSoft	AWS Marketplace listings show six-figure annual pricing tiers and services ecosystem.	Integration/control-plane products command enterprise budgets when operationally central.
Fiddler	Fiddler describes the AI control plane as standardized telemetry, reliable evaluation, continuous monitoring, enforceable policy, and auditable governance.	The market language is moving toward control-plane positioning for AI oversight.

12. Pricing Evidence and Commercial Model

The pricing model should separate discovery, pilot, annual license, implementation services, and channel economics. Mixing those together creates confusion because a paid pilot often includes substantial services, while an annual license is primarily software and support.

Recommended commercial architecture

Tier	Name	Target range	Duration	Primary deliverable
0	Discovery / Readiness	USD/EUR 5K-15K	1-2 weeks	Fit assessment, workflow candidate, deployment path, pilot scope.
1	Validation Sandbox	USD/EUR 15K-35K	30-45 days	Synthetic or representative workflow validation, starter kit, evidence sample.
2	Controlled Workflow Pilot	USD/EUR 45K-75K	45-90 days	One production-adjacent workflow, guided deployment, evidence package, support window.
3	Regulated Evidence Pilot	USD/EUR 75K-150K	60-120 days	More stakeholder review, stronger evidence package, broader operational validation.
4	Strategic Anchor Deployment	USD/EUR 150K-250K+	90-180 days	Multi-workflow or strategic partner deployment with extended implementation and roadmap input.

Annual licensing model

Scenario	Annual software range	Rationale
Conservative	USD/EUR 35K-75K	Anchored to lower-end AI governance, policy enforcement, and early controlled deployment.
Realistic	USD/EUR 75K-150K	Reflects runtime enforcement, evidence generation, customer-controlled deployment, and partner support.
Aggressive	USD/EUR 150K-300K+	Appropriate for regulated, multi-workflow, high-usage, or strategic deployments with broader operational value.

Pricing assumptions

Assumption	Commercial effect
Pilot pricing includes services	Discovery, mapping, deployment, validation, training, support, and evidence closeout are bundled into the pilot fee.
Annual license excludes one-time implementation	Annual licensing should not be judged against pilot fee directly; the pilot is services-heavy.
Partner-led delivery changes margin structure	Partner may earn services margin, referral fee, resale margin, or revenue share depending on role.
Configuration maturity increases pricing power	Policy builder, workflow wizard, and governance packs reduce implementation cost and support higher repeatability.
Security review increases enterprise readiness	Third-party security review is a future milestone that should improve procurement confidence.

13. Partner Economics and Channel Strategy

The channel model matters because CerbaSeal is likely to reach early EU or regulated customers through advisory and implementation partners. The partner is not merely a reseller; the partner may own client relationship, regulatory framing, implementation support, pilot scoping, and ongoing account management.

Model	When it fits	Revenue logic
Referral	Partner introduces customer but Lamont Labs sells and delivers.	Referral fee or commission; lowest partner workload.
Reseller	Partner sells CerbaSeal into client account with defined discount or margin.	Partner margin generally reflects sales ownership but not full implementation burden.
Implementation partner	Partner configures, deploys, trains, and supports customer.	Partner earns services margin; Lamont Labs earns software license and/or support fee.
Strategic EU channel	Partner has regional relationships, regulatory expertise, and possible exclusivity.	Requires written commitments, territory, duration, minimum activity, exit terms, and IP protection.

Public AWS Marketplace partner research reports professional-services gross margins of 25% to 45% and average software resale gross margins around 13%. That supports a channel structure where partners earn meaningful services revenue while Lamont Labs protects product IP and core licensing value.

Exclusivity considerations

Question	Why it matters
What triggers exclusivity?	Avoid granting exclusivity before a pilot, minimum activity, or signed client commitment.
What keeps exclusivity active?	Define minimum outreach, pilot targets, revenue milestones, or time-based review.
What happens if no pilot occurs?	Prevent indefinite lock-up without commercial progress.
What territory and market are covered?	EU-only language must be precise; consider sector boundaries if needed.
What rights remain with Lamont Labs?	Protect core IP, non-EU markets, inbound opportunities, and roadmap control.

Commercial posture: remain open to exclusivity discussion, but require written terms, clear commitments, limited duration, exit conditions, and protected IP ownership.

14. Adoption Roadmap and Build Priorities

The highest-leverage roadmap items are the ones that reduce implementation friction, partner dependency, and founder involvement without changing the core enforcement model. Additional invariants and more tests are useful, but the commercial bottleneck is repeatable delivery.

Priority	Build item	Why it matters	Expected commercial impact
1	Policy pack builder	Reduces the need for partners to hand-author JSON.	Moves configuration from developer work to guided governance setup.
2	Workflow mapping wizard	Turns discovery into a repeatable process.	Reduces consulting effort and founder support during first pilot.
3	Governance pack library	Provides templates for finance, insurance, KYC, claims, and AI agent workflows.	Reduces blank-page friction and improves buyer relevance.
4	Replay and evidence explorer	Makes audit and evidence review accessible to non-engineering stakeholders.	Improves compliance, partner, and executive understanding.
5	Partner deployment automation	Further reduces setup time and support load.	Raises partner readiness and lowers onboarding cost.
6	Security review package	Supports procurement, CISO review, and enterprise readiness.	Improves trust and accelerates review process.

Founder independence metrics

Metric	Current target direction
Typical founder hours per pilot	Move from founder-supported to partner-led, with Tier 3 reserved for core changes or unresolved defects.
Configuration burden	Move from JSON authoring by trained partners to guided policy builder and validation.
Deployment burden	Move from documented deployment to automated deployment with health checks and proof validation.
Evidence explanation burden	Move from founder-led explanation to visual evidence explorer and closeout template.
Partner readiness	Move from trained partner can deploy/configure to partner can lead pilots independently.

15. Risk Register and Mitigation Plan

The strongest external posture is neutral, factual, and operational. Risks should be described as current scope, deployment consideration, roadmap item, or operator responsibility rather than as deficiencies. The goal is to build trust through clarity.

Risk area	Neutral description	Mitigation
Security review	Formal third-party security review is planned as a future milestone.	Communicate transparently during pilot kickoff; prepare security review package.
HTTP access control	The starter server is intentionally minimal; network-level access control is implemented by operator.	Document gateway, VPN, private network, and deployment perimeter assumptions.
Identity provider integration	CerbaSeal evaluates TrustState but does not operate the client IdP.	Clarify integration boundary and required identity assertions.
Configuration complexity	Policy files provide flexibility but may require partner training.	Build policy pack builder, validation scripts, and workflow wizard.
Production monitoring	Operational monitoring is environment-specific.	Provide recommended signals, health endpoint, and integration notes.
Founder dependency	Early pilots may require Lamont Labs involvement for Tier 3 support and roadmap-level changes.	Use guided independence model, partner certification, and support boundaries.
Pricing uncertainty	Private enterprise pricing is not fully public across comparable vendors.	Use evidence-tiered pricing framework and revise after first pilots.

External communication posture

Use neutral language: current scope, roadmap item, operator responsibility, client-controlled deployment, guided implementation, defined escalation, and pilot evaluation. Reserve audit-style language such as critical gap, founder bottleneck, or hard blocker for internal analysis only.

16. Recommended Next Steps

The immediate objective is to prepare for a practical partner conversation: what can be piloted, what a pilot should cost, what a partner can deliver, what Lamont Labs retains, and what needs to be built next to make delivery repeatable.

Horizon	Recommended actions
Next 7 days	Finalize this dossier, prepare a 1-page commercial summary, create pricing appendix, and prepare concise discussion questions for Olivia and Tina.
Next 30 days	Build or improve policy pack builder specification, workflow mapping wizard, deployment automation, and pilot proposal template.
Next 60 days	Run internal dry-run of a 1-day pilot, measure founder involvement, create evidence closeout package, refine partner certification exercises.
Next 90 days	Prepare first paid pilot offer, define partner agreement structure, confirm IP language, and map pilot pricing to support hours and scope.
Before any exclusivity	Require written terms, minimum partner commitments, territory definition, exit conditions, IP protection, and compensation structure.

Recommended questions for the next partner call

Question	Why ask it
What does a successful first pilot look like from your side?	Defines success in partner language before pricing is locked.
What type of first client do you expect to approach?	Determines whether pilot should be simple validation or regulated evidence package.
What implementation role do you expect the partner to own?	Clarifies channel model and founder involvement.
What pricing model are you seeing in your market research?	Lets partner contribute market evidence and own local-market assumptions.
What commitments would come with any EU exclusivity?	Protects Lamont Labs from granting territory without reciprocal activity.

Overall assessment

CerbaSeal has moved beyond a pure technical proof. The credible near-term path is a controlled workflow pilot, priced as a guided governance validation engagement, delivered through a partner-supported operating model, and strengthened by productization work that reduces configuration and deployment friction.

17. Partner Discussion Brief

This page is designed for the immediate partner conversation. The goal is to give Olivia a clear view of what CerbaSeal is, what has been built, what a first pilot should validate, and what commercial questions need to be answered before a formal offer is prepared.

What Olivia needs to know

Question	Practical answer
What is CerbaSeal?	A deterministic governance enforcement layer that sits between a proposed decision and execution.
What does it produce?	ALLOW, HOLD, or REJECT decision plus reason codes, audit entries, and evidence bundle.
Is it a workflow system?	No. It integrates with existing workflow, identity, case, AI, and operational systems.
Can it be deployed by a client or partner?	Yes, using guided Docker Compose or Node.js Direct paths with configuration and verification steps.
What is the first pilot?	One governed workflow, one approval model, one deployment path, and one evidence package.
What should pricing reflect?	Discovery, workflow mapping, deployment support, validation, evidence review, support window, software value, and partner role.
What should not be promised?	Production certification, regulatory compliance certification, unlimited custom development, or 24/7 support without written scope.

Discussion priorities

Priority	Question to resolve
Pilot client profile	Which client type is most likely to buy first: seed/Series A AI company, fintech, insurance, regulated SaaS, or enterprise AI team?
Pilot scope	Which workflow has the cleanest path to evidence within 45-90 days?
Commercial structure	Will the partner charge the client directly and compensate Lamont Labs by percentage, fixed technical fee, or hybrid model?
Support window	How many partner/client support hours are included, how are they scheduled, and what triggers Tier 3 escalation?
Exclusivity	What concrete commitments, time limits, territory, exit terms, and minimum activity would justify exclusivity?

Best meeting objective: leave with enough agreement to draft a one-page heads-of-terms outline covering pilot scope, compensation structure, support boundaries, IP ownership, and exclusivity principles.

18. Pilot Commercial Structure

The first pilot should be priced as a guided governance validation engagement. It should not be described as a free trial, broad consulting project, or production deployment. The offer should clearly separate software value, implementation effort, partner involvement, and Lamont Labs support.

Recommended first pilot package

Component	Included
Discovery and readiness	Workflow candidate review, client readiness check, deployment mode selection, success criteria.
Workflow mapping	Actor roles, approval path, proposed actions, policy pack, GovernedRequest mapping.
Deployment support	Docker Compose or Node.js Direct setup, config/policy validation, health check, audit-chain verification.
Scenario validation	ALLOW, HOLD, and REJECT test cases against representative workflow.
Evidence closeout	Proof export, proof verification, evidence review session, pilot outcome summary.
Support window	Bounded support hours, partner-led Tier 1/2 support, Lamont Labs Tier 3 escalation for core issues.

Commercial model options

Model	How it works	When it fits
Fixed technical fee	Lamont Labs receives a fixed fee per pilot; partner prices client separately.	Best for early pilots when client pricing is uncertain but Lamont Labs support burden must be bounded.
Percentage split	Lamont Labs receives agreed percentage of client pilot fee.	Best when partner owns client relationship and pricing is transparent.
Hybrid	Base technical fee plus percentage upside.	Best for balancing support protection with upside from larger engagements.
License plus services	Client pays software license plus partner implementation services.	Best after first pilot once scope and repeatability are clearer.

Suggested negotiation posture

Topic	Recommended stance
Pricing	Do not name a final number until pilot scope, support expectations, and partner responsibilities are defined.
IP	Lamont Labs retains CerbaSeal core IP. Client-specific workflow artifacts and derivative improvements need written treatment.
Support	Bounded, scheduled, and role-defined. Avoid open-ended availability.
Exclusivity	Open to discussion only with written terms, activity commitments, limited duration, and exit conditions.
Pilot success	Define before work begins. Evidence package, operational sustainability, client value, and next-step clarity.

This structure gives the partner flexibility to test the market while protecting Lamont Labs from unpaid scope expansion, unlimited support, or premature exclusivity.

Appendix A - Reason Codes

Reason code	Category	Meaning
DECISION_ALLOWED	Terminal	All checks passed and authorization issued.
DECISION_HELD	Terminal	Request is valid but required approval is missing.
DECISION_REJECTED	Terminal	Invariant or policy rule failed.
NO_POLICY_PACK	INV-01	policyPackRef is null.
NO_PROVENANCE	INV-02	provenanceRef missing or incomplete.
REQUIRED_APPROVAL_MISSING	INV-03	Approval required but approvalArtifact absent.
INVALID_APPROVAL_AUTHORITY	INV-03	Approval request mismatch, timestamp issue, invalid approver, or chain mismatch.
PRIVILEGED_AUTH_NOT_SATISFIED	INV-03	Step-up or privileged authentication not satisfied.
APPROVAL_SIGNATURE_MISSING	INV-03	Approval signature field empty or whitespace.
LOGGING_NOT_READY	INV-04	loggingReady is false.
AI_CANNOT_AUTHORIZE	INV-05	AI proposal or AI actor attempts authority boundary violation.
PROHIBITED_USE	INV-10	prohibitedUse true.
CONTROL_STALE_OR_INVALID	INV-08	Sensitive controls stale or invalid.
TRUST_STATE_INVALID	INV-09	trustState.trusted false.
UNKNOWN_ACTION_CLASS	INV-11	Action class not registered.
MALFORMED_REQUEST	INV-11 / fallback	Schema validation or controlled fallback error.
INVALID_PROPOSAL	INV-12	proposedActionClass does not match proposal.requestedActionClass.

Appendix B - Deployment and Pilot Checklist

Step	Checklist item	Owner
1	Confirm pilot workflow, actor roles, approval path, and deployment mode.	Client + partner
2	Confirm hosting environment, network controls, persistence path, and operational owner.	Client
3	Prepare cerbaseal.config.json and cerbaseal.policy.json.	Partner + client review
4	Run installation path: Docker Compose or Node.js Direct.	Partner engineer
5	Run pnpm test and pnpm audit:repo.	Partner engineer
6	Run setup or verify config/policy files.	Partner engineer
7	Run ALLOW, HOLD, REJECT verification scenarios.	Partner + client
8	Confirm GET /health returns ok and auditChainValid true.	Partner engineer
9	Run representative pilot requests.	Client + partner
10	Export proof and verify proof.	Partner
11	Review evidence package with operational and compliance stakeholders.	Client + partner
12	Document issues, support tickets, and next-step recommendation.	Partner + Lamont Labs as needed

Partner readiness checklist

Readiness area	Partner should be able to demonstrate
Positioning	Explain CerbaSeal as governance enforcement infrastructure, not workflow software or compliance software.
Deployment	Stand up deployment starter and validate health without Lamont Labs intervention.
Configuration	Author a policy for a new workflow using actor mappings, approval chains, workflow rules, and action policies.
Verification	Run ALLOW, HOLD, and REJECT scenarios and interpret output.
Evidence	Export and verify proof package; explain evidence artifacts to a client stakeholder.
Support	Resolve common deployment and configuration issues using support guide before escalation.
Escalation discipline	Know which issues require Lamont Labs Tier 3 involvement.
Pilot leadership	Guide a client through discovery, mapping, deployment, verification, pilot run, evidence closeout, and handoff.

Appendix C - Source Notes

This section records the basis for major claims used in the dossier. Public market evidence is included as category and pricing context, not as a guarantee that CerbaSeal will command the same pricing.

Source	Evidence used
CerbaSeal-Core v0.1.0 System Architecture and Technical Reference Manual	Internal source manual: architecture, invariants, configuration, audit/evidence, API, deployment, tests, partner delivery, security.
AWS Marketplace - Styra Authorization Platform	Public marketplace pricing: Styra DAS Enterprise at \$5,000/month; runtime policy enforcement benchmark. URL: aws.amazon.com/marketplace/pp/prodview-4qemlqp6lodwg
AWS Marketplace - Workato Enterprise Automation Platform	Public marketplace pricing: Enterprise workspace with 1,000,000 tasks and Enterprise Support at \$143,750/year. URL: aws.amazon.com/marketplace/pp/prodview-yrn5rvs3n6v4c
AWS Marketplace - MuleSoft listings	Public marketplace evidence of six-figure middleware/integration pricing and professional-services ecosystem. URL: aws.amazon.com/marketplace
AWS Marketplace - Credo AI Enterprise AI Governance Platform	Private-offer pricing model based on contract duration, terms, and usage. URL: aws.amazon.com/marketplace/pp/prodview-x67krdatcdday
Fiddler AI Control Plane materials	AI control-plane positioning around enforceable policy, auditable governance, and centralized oversight. URL: fiddler.ai/control-plane
EU AI Act / Article 26 resources	Regulatory context for human oversight, monitoring, logging, and deployer responsibilities. URL: artificialintelligenceact.eu/article/26/
AWS Marketplace / Forrester partner economics summaries	Channel partner benchmarks: professional-services margins 25%-45%, resale margin around 13%. URL: aws.amazon.com/blogs/apn/a-total-economic-impact-partner-opportunity-analysis-for-channel-partners-selling-in-aws-marketplace/

Important limitations

Many enterprise AI governance, control-plane, and compliance vendors use private offers or custom pricing. Exact private contract values are often not publicly available. The pricing framework in this document should be used as a structured discussion model and refined after direct partner feedback, first pilot scoping, and observed support effort.

CerbaSeal-Core

Commercialization, Pilot Readiness & Market Positioning Dossier

Prepared for authorized Lamont Labs and partner evaluation. This document is intended to support technical understanding, pilot planning, pricing discussion, partner delivery design, and commercialization strategy. Redistribution should follow the applicable partner or confidentiality agreement.

Lamont Labs